# project\_report

June 27, 2016

## 1 Project 4 Report

This is my report of Project 4 of the Machine Learning Engineer Nanodegree - **Teach a Smartcab How to Drive**. In this project I use reinforcement learning techniques to teach an engine how to play a simple game of reaching a destination on a grid-like world given some restrictions. The project description can be found here, and my code for the project is on this github repo.

#### 1.1 Rules of the Game

Our agent should get to its destination on time and respect the following traffic rules:

- If the light is red, the agent generally cannot move but it can turn right when there is no traffic from the left going forward and no oncoming traffic turning left
- If the light is green, the agent generally can move but it cannot turn left when there is oncoming traffic going forward or turning right

At each moment in the game, the agent has 4 possible actions to choose from: it can go forward, turn right, turn left, or do nothing.

#### 1.2 Task 1: Implement a Basic Driving Agent

According to the project's instructions, the basic driving agent should "produce some random move/action." That's easy enough with something like this:

```
In [1]: import random

def update(self, t):
    # Gather inputs
    self.next_waypoint = self.planner.next_waypoint() # from route planner
    inputs = self.env.sense(self)
    deadline = self.env.get_deadline(self)

# Do something random
    action = random.choice(['right', 'left', 'forward', None])
```

We'll see below how this strategy fares in avoiding illegal moves and getting the agent to its destination. Spoiler: not very well.

## 1.3 Task 2: Identify and Update State

The following inputs are available for the agent at each update:

- *Light*: whether the light is red or green. As mentioned above, a green light means the agent can perform the next action, with the possible exception of a left turn, while a red light means the agent should stay put, with the possible exception of a right turn. So, it is important to add the light to the state.
- *Oncoming*: whether there is oncoming traffic, and which direction it is going. As mentioned above, oncoming traffic may mean the agent cannot turn left on a green light or right on a red light, so this information needs to be in the state as well.
- *Right*: whether there is traffic from the right of the agent, and which direction it is going. Right-of-way rules don't mention traffic to the right, so this is unnecessary information that doesn't need to be in the state for the agent to learn the optimal policy.
- *Left*: whether there is traffic from the left of the agent, and which direction it is going. Traffic from the left going forward means the agent cannot turn right on a red light, so this needs to be in the state.
- *Next waypoint*: the direction the agent should go to reach the destination. Without this information, the agent does not know where to go next and might as well wonder around randomly, so this needs to go in the state.
- Deadline: how much time the agent has left to reach its destination. At first, I would say this is not meaningful information for the agent, since it doesn't change right-of-way rules nor the best route. I thought about adding it to the state anyway, but this would mean a large increase in the number of possible states. Using only light ('red' or 'green'), oncoming (None, 'left', 'right', or 'forward'), left (None, 'left', 'right', or 'forward') and next\_waypoint ('left', 'right', or 'forward'), we have  $2 \times 4 \times 4 \times 3 = 96$  possible states. Adding deadline would mean multiplying this number by 50, if not more. So I'll keep deadline off my state for now.

So, the state will consist of a (light, oncoming, left, next\_waypoint) tuple.

I thought about combining inputs to come up with a state that would have information regarding whether or not it is okay to perform a given action. This would look like (ok\_forward, ok\_left, ok\_right, next\_waypoint) - a tuple with 4 items, but the number of overall states would be smaller:

## • State with raw inputs

- light can be 'green' or 'red'
- oncoming can be None, 'forward', 'left' or 'right'
- left can be None, 'forward', 'left' or 'right'
- next\_waypoint can be 'forward', 'left' or 'right'
- Full state space:  $2 \times 4 \times 4 \times 3 = 96$  possible states
- Preprocessed state ok\_forward can be True or False ok\_left can be True or False ok\_right can be True or False next\_waypoint can be 'forward', 'left' or 'right' Full state space:  $2 \times 2 \times 2 \times 3 = 24$  possible states

However, I checked on the Udacity forum and I'm not allowed to preproces inputs to come up with a simpler state space. (I did it anyway, but it's at the end of this report in a **Bonus** section.) So I'll implement the first type of state presented above, like this:

#### 1.4 Variables of Interest

Let's go a little deeper on what constitute a good agent in this setting. Some ideas come to mind:

- A good agent would reach its destination as fast as possible, and learn to do so well over time
- A good agent would make no mistakes
- A good agent would have a vast knowledge of the state space

So I'll keep track of the following variables:

- n\_dest\_reached: the number of times the agent has reached its destination
- last\_dest\_fail: the most recent trial in which the agent did not reach its destination
- sum\_time\_left: the sum over all trials of the steps the agent still had available when it reached its destination (0 if it never reached it)
- n\_penalties: the number of penalties incurred by the agent
- last\_penalty: the most recent trial in which the agent received a penalty
- len\_qvals: how many Q-values are mapped in the agent's Q-function that is, how many (state, action) pairs it visited during the simulation

I was also keeping track of the rewards accumulated by the agent, but decided to drop this information, because of two reasons:

- 1. The rewards are a way for the agent to do what's described above, not an end in and of themselves. If you're teaching a program to play chess or another board game, you can give it rewards for moves perceived as good, but at the end of the day what matters is not how many good moves the agent performs but how often it wins.
- 2. I will be analyzing the agent's performance over 100 trials. Since the penalty for an incorrect move is smaller than a reward for a correct one, an agent that sometimes perform an incorrect action, and because of that spends more moves getting to a destination, could end up with more rewards than an agent that only performs correct moves! (If I were to analyze the agent over a given number of *moves*, this would be different, because a more direct approach would lead to a larger number of destinations reached.)

## 1.5 Basic Agent Implementation

Putting together the update function, the state to implement and the variables of interest leads to this implementation of a basic agent:

```
In [3]: class BasicAgent():
            """A basic agent upon which to build learning agents."""
            def __init__(self, env):
                super(BasicAgent, self).__init__(env) # sets self.env = env, state
                self.color = 'red' # override color
                self.planner = RoutePlanner(self.env, self) # simple route planner
                self.qvals = {} # mapping (state, action) to q-values
                self.time = 0 # number of moves performed
                self.possible_actions = (None, 'forward', 'left', 'right')
                self.n_dest_reached = 0 # number of destinations reached
                self.last_dest_fail = 0 # last time agent failed to reach destinata
                self.sum_time_left = 0 # sum of time left upon reaching destination
                self.n_penalties = 0 # number of penalties incurred
                self.last_penalty = 0 # last trial in which the agent incurred in a
            def reset(self, destination=None):
                self.planner.route_to(destination)
            def best_action(self, state):
                Return a random action (other agents will have different policies)
                return random.choice(self.possible_actions)
            def update_qvals(self, state, action, reward):
                Keeps track of visited (state, action) pairs.
                (other agents will use reward to update
                the mapping from (state, action) pairs to q-values)
                self.qvals[(state, action)] = 0
            def update(self, t):
                # Gather inputs
                self.next_waypoint = self.planner.next_waypoint() # from route pla
                inputs = self.env.sense(self)
                deadline = self.env.get_deadline(self)
                # update time
                self.time += 1
                # Update state
                self.state = (inputs['light'], inputs['oncoming'], inputs['left'],
```

```
# Pick an action
action = self.best_action(self.state)

# Execute action and get reward
reward = self.env.act(self, action)
if reward < 0:
    self.n_penalties += 1

# Update the q-value of the (state, action) pair
self.update_qvals(self.state, action, reward)</pre>
```

A file with this version of the agent can be found here. (In the actual code, the implementation is class <code>BasicAgent(Agent)</code>.)

All other agents presented in this report will build upon this BasicAgent class, changing only the best\_action and update\_qvals methods. I altered some code in the simulator to update some BasicAgents values, such as n\_dest\_reached. My version of the simulator is here.

#### 1.5.1 Basic Agent Results

50%

Here are the results of running 100 simulations, of 100 trials each, with the basic agent presented above:

```
In [4]: from smartcab.run import run sims
        from smartcab.basic_agent import BasicAgent
        df_basic = run_sims(100, 100, BasicAgent)
        df_basic.to_csv('basic_agent_results.csv')
        df_basic.describe()
Out[4]:
               n_dest_reached last_dest_fail
                                               sum_time_left n_penalties
                  100.000000
                                   100.000000
                                                  100.000000 100.000000
        count
                    20.110000
                                    99.750000
                                                  301.090000 1588.810000
       mean
        std
                     4.242391
                                     0.575159
                                                   76.095381
                                                                66.093245
       min
                     8.000000
                                    97.000000
                                                  125.000000 1402.000000
        25%
                    18.000000
                                   100.000000
                                                  253.250000 1546.750000
        50%
                    20.000000
                                   100.000000
                                                  306.000000 1590.000000
        75%
                    23.000000
                                   100.000000
                                                  345.250000 1631.500000
                                                  480.000000 1788.000000
       max
                    30.000000
                                   100.000000
               last_penalty
                              len_qvals
                      100.0 100.000000
        count
       mean
                      100.0
                             86.020000
        std
                              6.747959
                        0.0
       min
                      100.0
                              69.000000
        25%
                      100.0
                              82.000000
```

86.000000

100.0

```
75% 100.0 91.000000 max 100.0 104.000000
```

A csv file with the results of the simulation above can be found here. The code used to run simulations on this and other agents below is here.

As expected, an agent that takes actions at random is a lousy one. I'm actually surprised it gets to its destination some 20 times on average. The number of penalties is quite large - around 16 per trial on average.

Another think to keep in mind: on average, a completely random agent will explore some 86 (state, action) pairs in 100 trials. We have  $96 \times 4 = 384$  such pairs (the 96 possible states times the 4 possible actions). So randomly picking actions in a state means exploring, on average, about 22% of (state, action) pairs.

This number seems low, but it may be due to a dearth of other cars in the smartcab's world. That makes it unlikely that any state with more than two cars in a given intersection will happen, for instance. The planner may also rely more heavily or certain actions, so that (state, action) pairs involving less used actions will be rarer.

#### 1.5.2 Perfect Agent Results

We can contrast the results of the basic agent with those of an agent that always picks the correct action, given by the best\_action method defined below:

```
In [5]: def best_action(self, state):
            Returns the best possible action.
            # retrieve state information
            light, oncoming, left, waypoint = state
            # retrieve best action
            action = waypoint
            # On a red light, the agent can only turn right, and even so only if:
            # - no oncoming traffic is going left
            # - no traffic from the left is going forward
            if light == 'red':
                if any([action != 'right', oncoming == 'left',
                        left == 'forward']):
                    action = None
            # On a green light, the agent cannot turn left if there is
            # oncoming traffic going forward or right
            elif action == 'left' and (oncoming == 'forward' or oncoming == 'right'
                action = None
```

The perfect agent implementation is here. These are the results of the perfect agent:

return action

In [6]: from smartcab.perfect\_agent import PerfectAgent

```
df_perfect = run_sims(100, 100, PerfectAgent)
df_perfect.to_csv('perfect_agent_results.csv')
df_perfect.describe()
```

Out[6]:	n_dest_reached	last_dest_fail	sum_time_left	n_penalties	\
count	100.000000	100.00000	100.000000	100.0	
mean	99.930000	3.39000	1794.330000	0.0	
std	0.293189	15.45204	79.488142	0.0	
min	98.000000	0.00000	1588.000000	0.0	
25%	100.000000	0.00000	1741.750000	0.0	
50%	100.000000	0.00000	1800.000000	0.0	
75%	100.000000	0.00000	1843.250000	0.0	
max	100.000000	100.00000	1997.000000	0.0	

	last_penalty	len_qvals
count	100.0	100.000000
mean	0.0	23.410000
std	0.0	2.796444
min	0.0	16.000000
25%	0.0	21.750000
50%	0.0	23.500000
75%	0.0	25.000000
max	0.0	31.000000

A csv file with the results of the simulation above can be found here.

Surprisingly, there are occasions in which this perfect agent does not reach its destination. This may be due to issues with the planner or with bad luck (getting a lot of red lights, for instance). On the other hand, the agent incurs no penalties whatsoever, which is to be expected.

Also note how fewer (state, action) pairs are explored. For the perfect agent, a state will always be paired with the same action (give by next\_waypoint), so there are only 96 possibilities for it to explore.

## 1.6 Task 3: Implement Q-Learning

Let's return to the basic agent and modify it to learn from its actions.

#### 1.6.1 Deciding on the Appropriate Q-learning Function

The general form of the *Q*-function is:

$$Q(s, a) = R(s) + \gamma \sum_{s'} T(s, a, s') \max_{a'} Q(s', a')$$

That is, the Q-value for a given (state, action) pair is the the reward for that state, R(s), plus the discounted expected value of Q for the next state the agent lands in, considering the transition function  $T(s,a,s') = \Pr(s' \mid s,a)$  (the probability of landing on state s' coming from state s and performing action s) and that, whatever s' is, the agent will pick s' so as to maximize s'0 from there on.

The *Q*-learning update function is given by:

$$\hat{Q}_t(s, a) = (1 - \alpha_t)\hat{Q}_{t-1}(s, a) + \alpha_t(r + \gamma \max_{a'} \hat{Q}_{t-1}(s', a'))$$

That is, our estimate of the Q-value for the (state, action) pair is updated with the learning rate ( $\alpha_t$ , which varies over time) by the observed reward (r) and our previous estimate of the Q-value for the observed next state (s'), discounted by the discount factor ( $\gamma$ ) and considering the agent will pick the action a' that maximizes Q from the next state on.

However, in this case there's no need to worry about the future state, since the agent gets an immediate reward for doing the right thing. According to the project description (emphasis added):

The smartcab gets a reward for each successfully completed trip. A trip is considered "successfully completed" if the passenger is dropped off at the desired destination (some intersection) within a pre-specified time bound (computed with a route plan).

It also gets a smaller reward for each correct move executed at an intersection. It gets a small penalty for an incorrect move, and a larger penalty for violating traffic rules and/or causing an accident.

So, even though the larger reward is only reaped once the agent reaches its destination, there are smaller rewards for following the correct path, and penalties for not doing so. This should be enough for the agent to learn the best policy.

Granted, ignoring the agent's future decisions means I'm not using some information that could be of help. But the upside is a simplification of the problem: it's as if the agent were playing a 1-round game over and over, with immediate rewards for immediate actions. I expect this simplification more than compensates ignoring long-term rewards in this particular setting.

#### 1.6.2 Q-learning implementation

This means I won't actually bother with keeping track of the state the agent ends up in after performing an action (or, to be more technical, I'm setting the discount factor  $\gamma$  to zero). My update function will then simply be  $\hat{Q}_t(s,a) = (1-\alpha_t)\hat{Q}_{t-1}(s,a) + \alpha_t r$ .

Here are the best\_action and update\_qvals methods that define such an agent:

```
def update_qvals(self, state, action, reward):
    """

    Updates the q-value associated with the (state, action) pair
    """

# define the learning rate for the current time
learn_rate = 1.0 / self.time

# update the q-value for the (state, action) pair
self.qvals[(self.state, action)] = \
    (1 - learn_rate) * self.qvals.get((self.state, action), 0) + \
    learn_rate * reward
```

This implementation can be found here.

#### 1.6.3 Learning Agent Results

Here are the results for this learning agent:

```
In [8]: from smartcab.learning_agent import LearningAgent
        df_learning = run_sims(100, 100, LearningAgent)
        df_learning.to_csv('learning_agent_results.csv')
        df learning.describe()
Out[8]:
                n_dest_reached
                                 last_dest_fail
                                                  sum_time_left
                                                                  n_penalties
                                                                   100.000000
        count
                     100.00000
                                     100.000000
                                                     100.000000
                      99.10000
                                      18.810000
                                                    1751.910000
                                                                    32.220000
        mean
        std
                       1.04929
                                      29.609136
                                                      69.277818
                                                                    14.602878
        min
                      93.00000
                                       0.000000
                                                    1555.000000
                                                                    19.000000
        25%
                      99.00000
                                       0.000000
                                                    1715.500000
                                                                    27.000000
        50%
                      99.00000
                                       1.000000
                                                    1746.000000
                                                                    30.000000
        75%
                     100.00000
                                      30.250000
                                                    1798.500000
                                                                    34.000000
                     100.00000
                                      99.000000
                                                    1899.000000
                                                                   160.000000
        max
                last_penalty
                                len_qvals
        count
                  100.000000
                              100.000000
                                51.370000
        mean
                   93.590000
        std
                    6.612041
                                 6.351099
        min
                   73.000000
                                36.000000
        2.5%
                   90.000000
                                47.000000
        50%
                   96.000000
                                51.000000
        75%
                   99.000000
                                56.000000
                                66.000000
                  100.000000
        max
```

A csv file with the results of the simulation above can be found here.

These are really good results! The average and median number of destinations reached are actually comparable to the perfect agent's, although the standard deviation is much higher (see above). Also note that when it comes to exploring the state space this agent's numbers are between the perfect agent's (that only explore the best actions provided by the planner) and the

basic agent (that does nothing but explore at random), with an average of 51 (state, action) pairs explored.

If there's room for improvement in this agent, it's in the last\_penalty variable: in over 75% of the simulations, the agent incurred in at least one penalty in the last 11 trials. On the bright side, 75% of simulations had 34 penalties or fewer - thats 0.34 penalties per trial, or 1 penalty for every 3 trials, which seems quite good specially considering there must be more penalties during the early trials, when the agent knows next to nothing about the world.

## 1.7 Task 4: Enhance the Driving Agent

The last task of the project, according to the project description, is:

Apply the reinforcement learning techniques you have learnt, and tweak the parameters (e.g. learning rate, discount factor, action selection method, etc.), to improve the performance of your agent. Your goal is to get it to a point so that within 100 trials, the agent is able to learn a feasible policy - i.e. reach the destination within the allotted time, with net reward remaining positive.

This goal has already been reached with the basic learning agent above. But there's still room for improvement, particulary regarding penalties late in the simulation.

According to one of the videos in the Reinforcement Learning course, there are three characteristics that can modify a Q-learning algorithm:

- How Q-values are initialized
- How the learning rate decays
- How the action is picked

I'll tackle these one at a time and discuss the results below.

#### 1.7.1 Changing Initial Q-Values

One Q-learning implementation briefly discussed in the Reinforcement Learning lessons for this project is "optimism in the face of uncertainty". The idea is that high initial Q-values (implying an "optimistic" agent in the sense that it initially believes all possible actions will yield excellent rewards) lead to an exploratory-leaning agent, because it will delay exploiting familiar paths, since those will end up with lower Q-values than its initial estimate.

All it takes for the learning agent above to become optimistic is changing the value it gets when the (state, action) pair is not yet a key in qvals (that is, when the pair is seen for the first time):

```
In [9]: def best_action(self, state):
    """

    Returns the best action (the one with the maximum Q-value)
    or one of the best actions, given a state, being
    optimistic in the face of uncertainty.
    """

# get all possible q-values for the state
# (be optimistic in the face of uncertainty)
all_qvals = {action: self.qvals.get((state, action), 100)
```

Instead of getting a 0 for previously unknown (state, action) pairs, the agent now gets a breathtaking 100, which is way more than even getting to the destination. So, it has more of an incentive to investigate new paths. (The implementation is here.)

I thought about also modifying the update\_qvals method to incorporate the initial optimistic Q-value to the updates, but that's unnecessary, since the point of optimistic initialization is to get the agent to test a given action in a given state. Once it has tested it, there's no need to use that initial optimistic value for anything. (It might be a different story if I were taking future states and actions into account, since dropping the value of a (state, action) pair might then mean very shallow observations of many possibilities.)

These are the results for running 100 simulations of 100 trials each with this optimistic agent:

```
In [10]: from smartcab.optimistic_agent import OptimisticAgent
         df_optimistic = run_sims(100, 100, OptimisticAgent)
         df_optimistic.to_csv('optimistic_agent_results.csv')
         df_optimistic.describe()
                n_dest_reached
Out [10]:
                                 last_dest_fail
                                                  sum_time_left
                                                                  n_penalties
                                      100.000000
                                                      100.000000
         count
                      100.00000
                                                                   100.000000
                       98.63000
                                       21.870000
                                                     1733.180000
                                                                     42.130000
         mean
                                       31.727901
                                                       95.359753
                                                                    17.743205
         std
                        1.70948
                       87.00000
                                                                     28.000000
         min
                                        0.000000
                                                     1293.000000
         25%
                       98.00000
                                        0.000000
                                                     1691.750000
                                                                     37.000000
         50%
                       99.00000
                                        1.000000
                                                     1734.000000
                                                                     39.000000
         75%
                      100.00000
                                       35.750000
                                                     1793.250000
                                                                     42.000000
         max
                      100.00000
                                      100.000000
                                                     1910.000000
                                                                   158.000000
                                len_qvals
                last_penalty
                    100.00000
                               100.000000
         count
                     94.26000
                                65.010000
         mean
         std
                      5.69178
                                  6.170138
         min
                     73.00000
                                48.000000
         25%
                     92.00000
                                61.000000
                     95.50000
         50%
                                65.000000
         75%
                     98.25000
                                69.000000
```

A csv file with the results of the simulation above can be found here.

100.00000

max

The results of the optimistic agent are comparable to those of the original learning agent. The main difference is that the number of penalties is somewhat larger, as well as the number of (state, action) pairs explored - both of which makes sense: optimism means more

83.000000

exploration, and more exploration in this setting means more penalties. When it comes to last\_penalty, however, optimistic initialization does not yield better results than the previous implementation.

#### 1.7.2 Modifying the Learning Rate Decay

The learning rate of my original learning agent is the inverse of time: when time is 1, the learning rate is 1; when time is 2, the learning rate is 1/2, and so on to 1/3, 1/4, 1/5 etc.

This drop seems very steep, but it can be modified. I decided to use the following equation to do so:

```
In [11]: def learn_rate(mult, time):
    return 1.0 / (1 + mult * time)
```

That means the learning rate will begin at 1 when time is zero, but its decrease will be a function of both time and the mult factor. If mult = 1, the learning rate is the same as above (assuming time starts at zero).

Since I think the drop in the learning rate is steep enough when  $\mathtt{mult} = 1$ , I won't bother with larger values of it, which would only magnify the drop. Instead, I'll work with 10 values of  $\mathtt{mult}$  from 1 to 0.001.

My prior on this experiment is that the learning rate will not have a significant impact on the learner. Because of how the rewards are structured (immediate negative rewards for incorrect or illegal moves, immediate positive rewards for the correct legal move, zero for staying put), and since only one action is correct at any one point, the rewards very quickly should take the following pattern:

- incorrect/illegal moves: minus something
- staying put: zero
- correct legal move: plus something else

The exact something and something else do not really matter: the correct legal move will always win out.

The only way I can think for something to go wrong is if an agent performs an incorrect move that *immediately takes it to the destination*. In this scenario, the positive reward for reaching the destination would overwhelm the negative reward for making an incorrect move, and, if the learning rate drops too quickly, no future penalties would be able to revert this. But, considering the simulation does not allow illegal moves (the agent takes the penalty for trying something illegal, but stays put), how would a legal move that takes the agent to the destination be an incorrect one?

Anyway, my point that I think tweaking with the learning rate will not be of much use given the rewards, the behavior of the agent and the states I came up with. Let's see if I'm right.

This is the update\_qvals() function for the learning agent with different learning rates (best\_action() stays the same):

```
In [12]: def update_qvals(self, state, action, reward):
    """

    Updates the q-value associated with the (state, action) pair
    """

# define the learning rate for the current time
learn_rate = 1.0 / (1 + self.mult*self.time)
```

```
self.qvals[(self.state, action)] = \
   (1 - learn_rate) * self.qvals.get((self.state, action), 0) + \
   learn_rate * reward
```

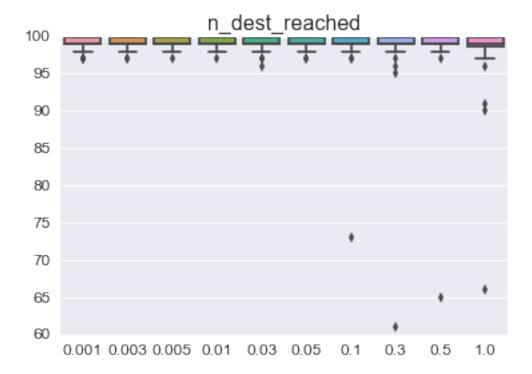
And this is the way I came up with to run the code using different values for mult:

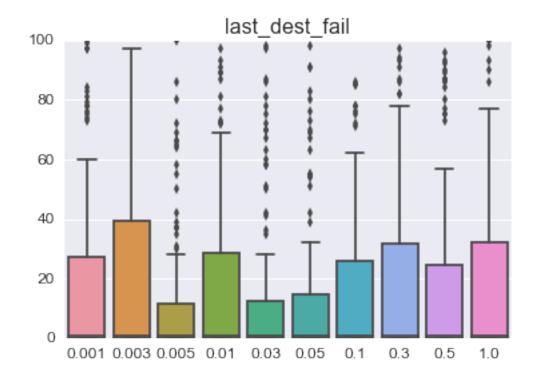
```
In [13]: def several_rate_changes(mults):
             For each mult value in mults, runs a simulation
             with a RateChangeAgent.
             Returns a dict with dataframe results for the agent
             for each mult value.
             results = {}
             for mult in mults:
                 mult_results = []
                 for i in range (100):
                     sim_results = run_rate_change(mult)
                     mult_results.append(sim_results)
                 df_results = pd.DataFrame(mult_results)
                 df_results.columns = ['n_dest_reached', 'last_dest_fail',
                                        'sum_time_left', 'n_penalties',
                                        'last_penalty', 'len_qvals']
                 df_results.to_csv("rate_change_{}_results.csv".format(mult))
                 results[mult] = df_results
             return results
```

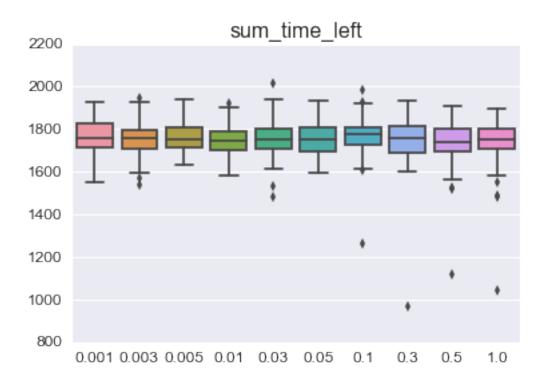
The code for studying changes in learning rates is here.

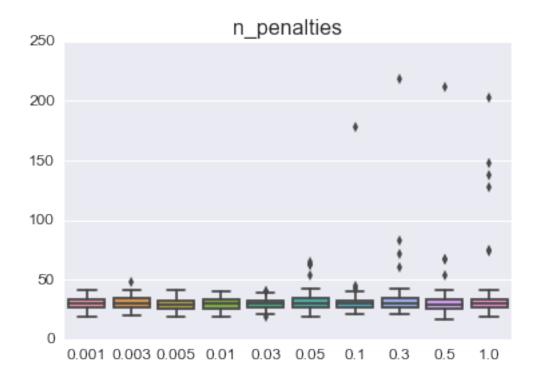
Let's see the results for these different learning agents with different learning rates - this time using plots to make it easier to compare values.

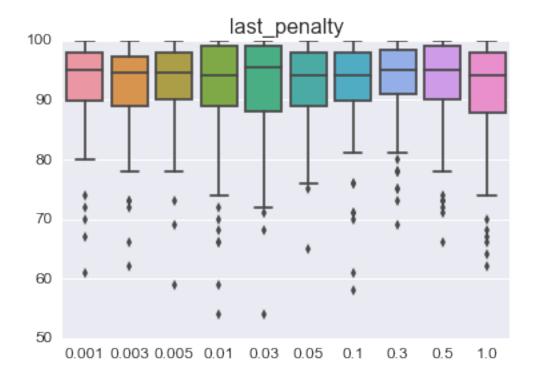
```
rate_panel = pd.Panel.from_dict(learning_rate_results)
for column in df_learning.columns:
    sns.boxplot(data=rate_panel.minor_xs(column))
    plt.title(column)
    plt.show()
    plt.close()
```

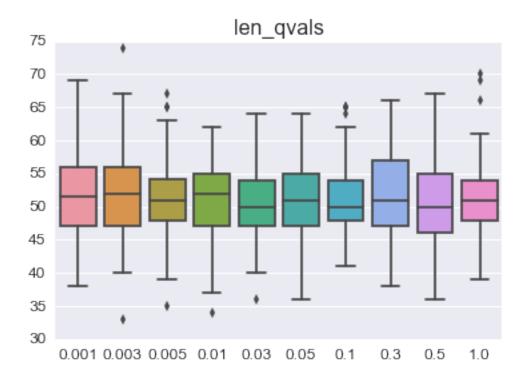












My posterior is very close to my prior: for this agent, tweaking the learning rate doesn't do much. last\_penalty, in particular, remains (on average) as high as it was before.

### 1.8 Randomly Picking Actions

Another possibility to tweak the learning agent is making it pick actions at random from time to time. The following best\_action() function does just that:

```
In [17]: def best_action(self, state):
             m m m
             Returns the best action (the one with the maximum Q-value)
             or a random action
             # get the rate of random values at this point in time
             random_rate = 1.0 - self.time * self.eps
             # if random number smaller than random rate,
             # the agent picks an unexplored action at the current state
             if random.random() < random_rate:</pre>
                 unexplored_actions = [action for action in self.possible_actions
                                        if (state, action) not in self.qvals.keys();
                 if unexplored_actions:
                     actions = unexplored_actions
                 else: # if no actions are unexplored in this state, pick any action
                     actions = self.possible_actions
```

(When picking at random, the agent will pick, if possible, an action that hasn't been performed yet given the state. This increases the exploratory nature of the agent somewhat.)

The larger random\_rate, the more probable it is that the agent will pick an action at random. eps is the variable that governs how fast this probability drops, much like mult governed how fast the learning rate would drop in the implementation above.

Note that at some point self.time \* self.eps will be larger than 1, and random\_rate will drop below zero, meaning it will be impossible for the agent to pick an action at random. The idea is that at some point the agent will end its exploratory activity completely, and dedicate itself to exploiting what it's learned. It is possible to modify the formula for random\_rate so that a small chance of randomness remains (using something like the formula for learn\_rate in the previous implementation), but in a setting such as teaching a smartcab to drive at some point you want to completely rule out random actions that translate to a larger chance of performing an illegal move or even causing an accident.

I ran several simulations with different values for eps using pretty much the same code as the one to vary learning rates:

```
In [18]: def several_random_changes(epses):
             For each eps value in epses, runs a simulation
             with a LearningRandomAgent.
             Returns a dict with dataframe results for the agent
             for each eps value.
             results = {}
             for eps in epses:
                 eps_results = []
                 for i in range (100):
                     sim_results = run_random_change(eps)
                     eps_results.append(sim_results)
                 df_results = pd.DataFrame(eps_results)
                 df_results.columns = ['n_dest_reached', 'last_dest_fail',
                                        'sum_time_left', 'n_penalties',
                                        'last_penalty', 'len_qvals']
                 df_results.to_csv("random_rate_{}_results.csv".format(eps))
                 results[eps] = df_results
             return results
```

The code for learning agents with probabilistically randomized behavior is here. And here are the results:

```
In [19]: from smartcab.learning_random_agent import several_random_changes
         epses = [1, 0.5, 0.3, 0.1, 0.05, 0.03, 0.01, 0.005, 0.003, 0.001]
         learning_random_results = several_random_changes(epses)
In [20]: import pandas as pd
         random_panel = pd.Panel.from_dict(learning_random_results)
         for column in df_learning.columns:
             sns.boxplot(data=random_panel.minor_xs(column))
             plt.title(column)
             plt.show()
             plt.close()
                             n dest reached
        100
         95
         90
         85
         80
         75
         70
         65
         60
```

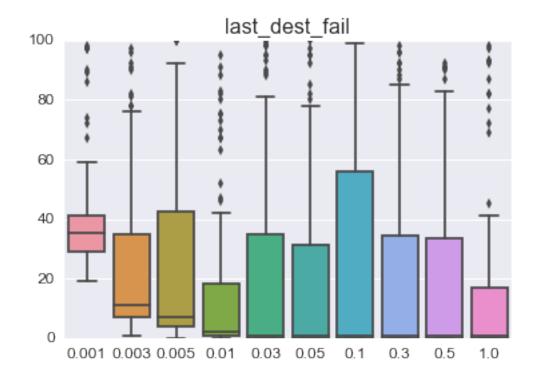
0.1

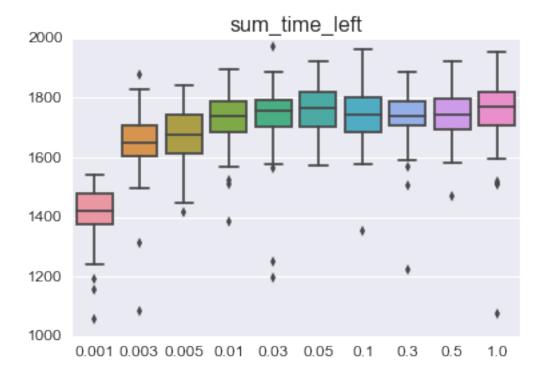
0.3

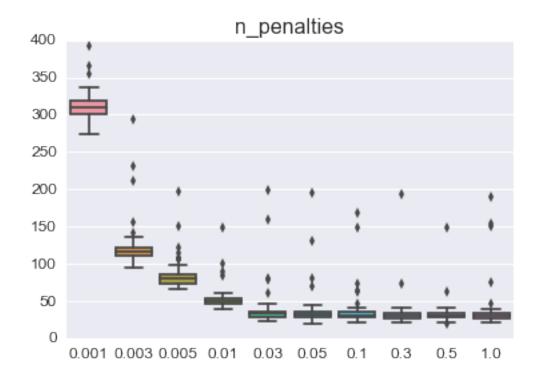
0.5

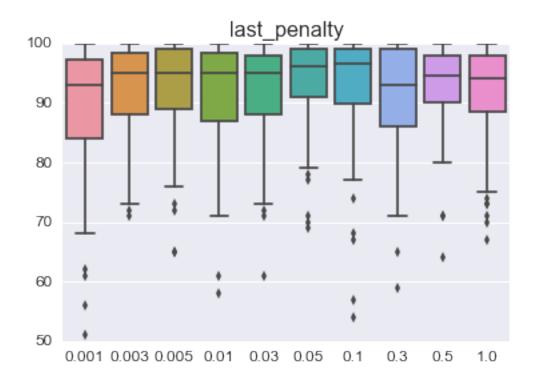
1.0

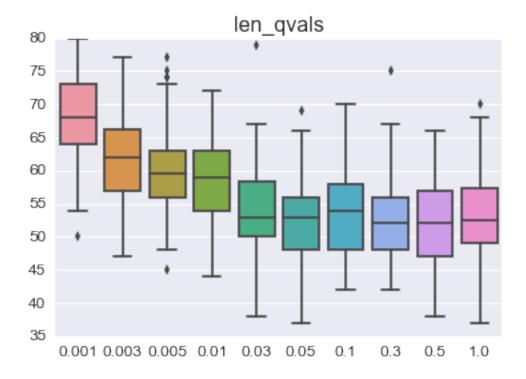
0.001 0.003 0.005 0.01 0.03 0.05











Agents with a smaller eps (that is, a larger propensity to explore at random) perform clearly worse in terms of destinations reached and number of penalties. They do, however, tend to explore more of the state space. Which suggests that maybe the exploration-exploitation dilemma is not that large of an issue in this setting: the original learning agent is not very exploratory, but exploring more does not have a clear upside.

#### 1.9 Other Ideas

No modification of the original learning agent is clearly better than it - and in fact it seems like, on the contrary, my original agent is better than all of the other possibilities analyzed. Given that the original learning agent had results comparable to those of a perfect agent, this is not very surprising.

However, there may be some room for improvement in other areas:

- Improve the planner. Even a perfect agent, defined as one that follows the planner's recommendation whenever possible, sometimes fails to reach its destination. A better planner could mean less mistakes for the agents.
- Increase the number of trials. Even when the agent picks actions completely at random, a large number of (state, action) pairs remains unexplored, suggesting some states may come up only rarely. With a larger number of trials, more of the state space can be explored. This could solve the problem of the learning agent still receiving penalties even in the last trials.
- Increase the number of other cars in the simulation. Maybe one reason why some states are so rare is the small number of cars on the streets. With more cars, it would be more likely to see full intersections, for instance. As it is now, I believe most of the time there are only one or two cars per intersection, and some states come up only rarely.

## 1.10 Project Rubric Rundown

Let's quickly review the project rubric and fill any remaining gaps:

- Implement a basic driving agent
  - Agent accepts inputs
  - Produces a valid output
  - Runs in simulator
- Identify and update state
  - Reasonable states identified
  - Agent updates state

The tasks above were implemented in the basic agent (code here). The states were identified and discussed in the section "Task 2: Identify and Update State" above.

- Implement Q-Learning
  - Agent updates Q-values
  - Picks the best action
  - Given the current set of Q-values for a state, it picks the best available action.
  - Changes in behavior explained

The tasks above were implemented in the learning agent (code here). The changes between agentes were presented in tables along the report. To summarize, the basic agent merely drives at random, and only reaches its destination by accident; while the learning agent presents a similar erratic beahvior at the beginning of the simulation but quickly starts aiming for the destination, indicating it is learning to pick as action the next waypoint. It also learns to avoid illegal moves as time goes by.

- Enhance the driving agent
  - Agent learns a feasible policy within 100 trials
  - Improvements reported
  - Final agent performance discussed

The original learning agent from the previous step was already able to learn a feasible policy within 100 trials, getting results that were close to a perfect agent's, as discussed above. I tried to tweak some parameters to improve the main problem with my learning agent (frequently getting late penalties), but none of the changes implemented were of help. These modifications include optimism in the face of uncertainty, modified learning rates and occasional random actions.

My final agent (code here) is actually the first agent I came up with. It gets very close to the perfect agent in terms of number of reaching its destination, but incurs in much more penalties, and occasionally is penalized in late trials. I believe this could be fixed by either increasing the number of trials (i.e., extending the learning period) or the number of cars in the world, so that more (state, action) pairs would be visited by the agent.

As stated and above and coded here, the perfect agent is one that always picks next\_waypoint as its action and obey traffic rules - that is, it's an agent whose action is always either next\_waypoint or, when said action cannot be performed, None.

## 1.11 Bonus 1: Changing the State Space

As mentioned above, I checked on the Udacity forum and I'm not allowed to do this. But what if I were to work on the inputs before passing them to the Q-learning function?

```
In [21]: def update(self, t):
             # Gather inputs
             self.next_waypoint = self.planner.next_waypoint() # from route planner
             inputs = self.env.sense(self)
             deadline = self.env.get_deadline(self)
             # update time and learning rate
             self.time += 1
             ok_forward = (inputs['light'] == 'green')
             ok_right = (inputs['light'] == 'green') or \
                 ((inputs['oncoming'] != 'left') and (inputs['left'] != 'forward'))
             ok_left = all([inputs['light'] == 'green',
                            inputs['oncoming'] != 'forward',
                            inputs['oncoming'] != 'right'])
             # Update state
             self.state = (ok_forward, ok_right, ok_left, self.next_waypoint)
             # Pick the best known action
             action = self.best action(self.state)
             # Execute action and get reward
             reward = self.env.act(self, action)
             if reward < 0:</pre>
                 self.n_penalties += 1
             # Update the q-value of the (state, action) pair
             self.update_qvals(self.state, action, reward)
```

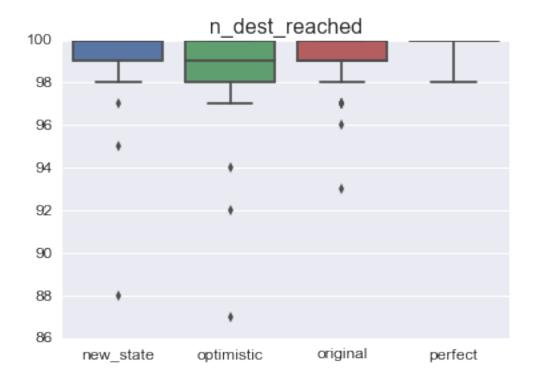
The code for this agent is here.

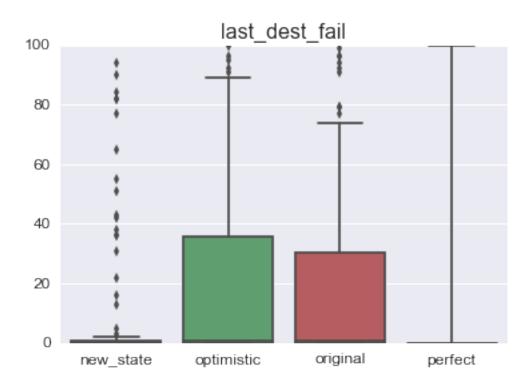
Now each state is a tuple with 4 values: whether it's okay to go forward, turn right, or turn left (2 options each), and the next waypoint (3 possibilities). Multiplying this by 4 possible actions, we have a slim 96 possible (state, action) pairs.

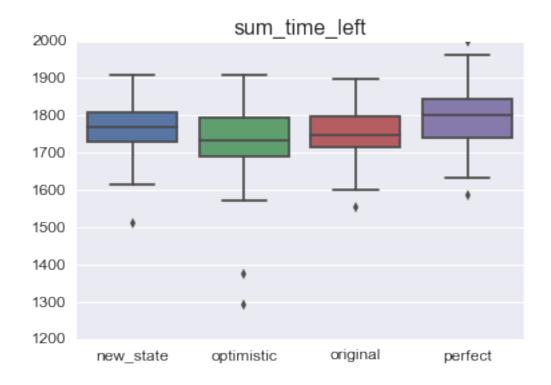
```
In [22]: from smartcab.new_state_agent import NewStateAgent
        df_new_state = run_sims(100, 100, NewStateAgent)
        df_new_state.describe()
Out [22]:
               n_dest_reached last_dest_fail sum_time_left n_penalties
                   100.000000
                                   100.000000
                                                 100.000000 100.000000
        count
                    99.310000
                                     9.920000
                                                 1766.720000
                                                              18.960000
        mean
                                   23.129709
                                                   72.239791
                                                              17.642877
        std
                     1.397653
```

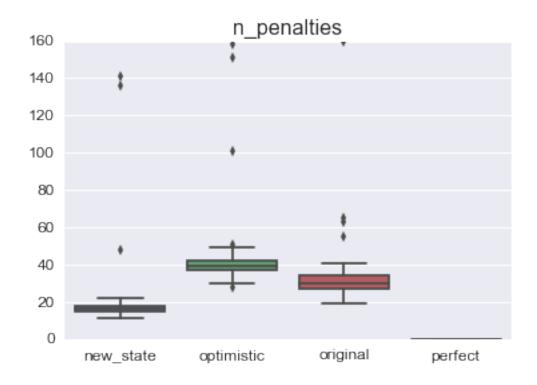
```
min
            88.000000
                               0.000000
                                           1512.000000
                                                           11.000000
25%
            99.000000
                               0.000000
                                           1730.250000
                                                           15.000000
50%
                              0.000000
                                           1767.500000
           100.000000
                                                           16.000000
75%
           100.000000
                              1.000000
                                           1809.500000
                                                           18.000000
                                           1910.000000
           100.000000
                              94.000000
                                                          141.000000
max
                       len_qvals
       last_penalty
         100.000000
                      100.000000
count
          71.150000
                       27.060000
mean
          25.120467
                        3.595227
std
           4.000000
                       20.000000
min
25%
          56.000000
                       25.000000
50%
          78.500000
                       27.000000
75%
          93.000000
                       30.000000
         100.000000
                       35.000000
max
```

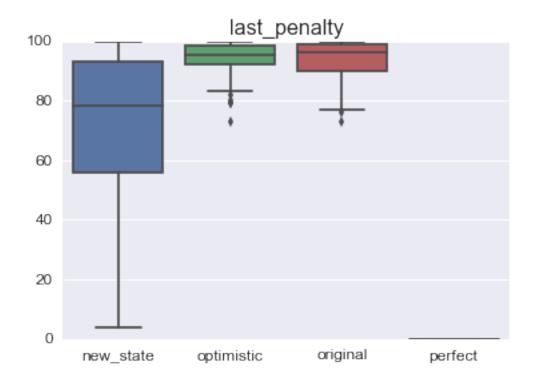
Alright, now last\_penalty is significantly down, and the number of destinations reached remains similar to the original learning agent. Let's visualize how this implementation compares to the original learning agent, the optimistic one and the perfect one:

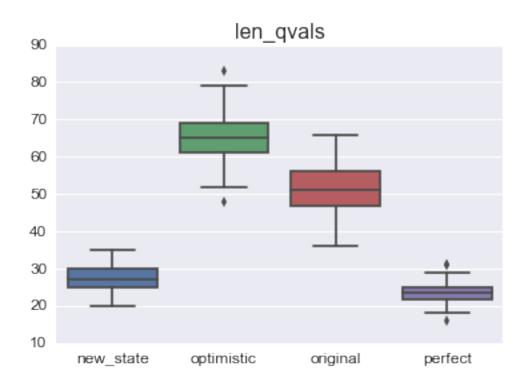










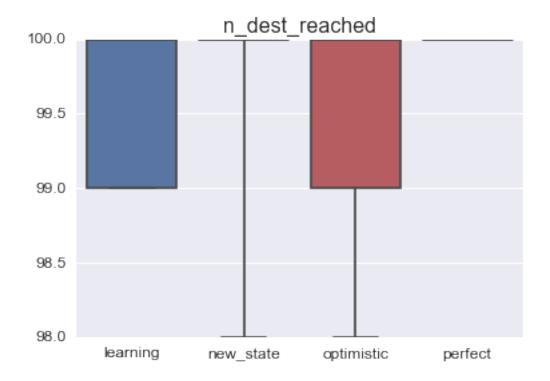


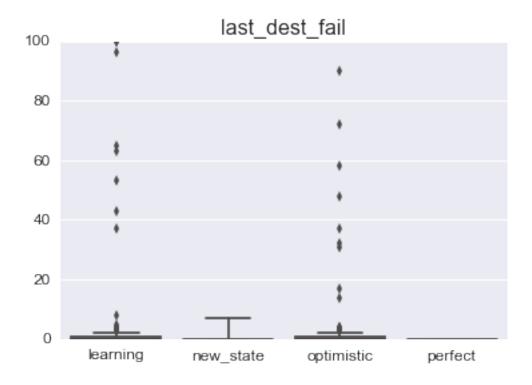
That's a clear improvement across the board, with a far simpler state space.

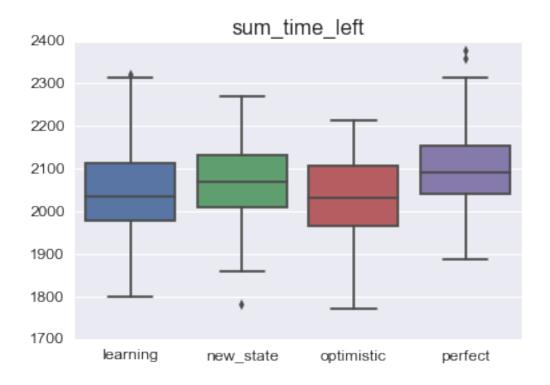
Is this cheating? Maybe. But imagine you're trying to use reinforcement learning to teach a program to play chess. I think it's fair to say you would hardcode how the pieces move and then let the learner play, instead of making it first figure out that pawns can't move back and bishops only move diagonally. So, it's reasonable to expect the rules of the game will be known in advance - and I'm not even doing that: I'm just passing a simpler state for the agent's Q-learning function, but it still needs to figure out how this state applies. So this might be fair game after all.

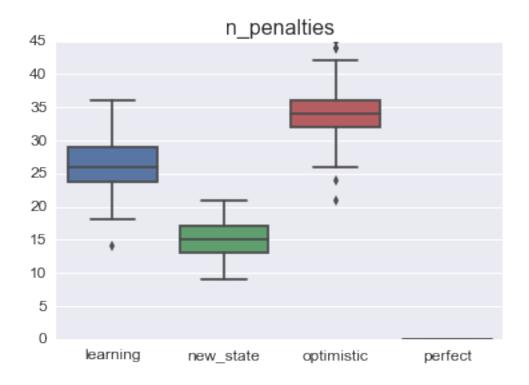
## 1.12 Bonus 2: Improving the Planner

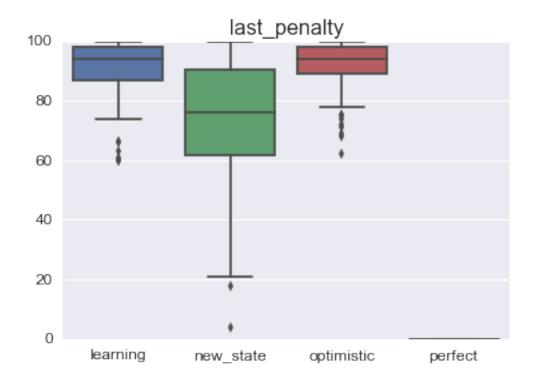
Let's see how the four main agents presented above - the original learning agent, the optimistic agent, the agent with a simpler state space, and the perfect agent - fare when taking waypoints from an improved planner I came up with.

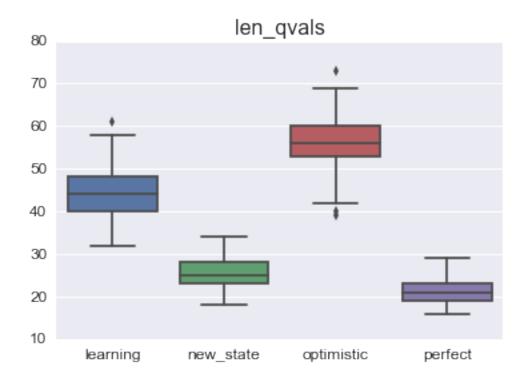












Now the perfect agent reaches its destination every time - but the other agents' performance also improved, and none of them fail to reach their destinations more than twice in any given simulation. And these rare failures tend to occur earlier than they did with the original planner.

The number of penalties - or at least the number of simulations with large number of penalties - also decreases, but the problem of late last penalties persists. As mentioned above, an environment with more cars may help with this, by forcing the smartcab to explore more of the state space in the same number of trials. Or maybe there could be a way of adapting the agent's best\_action method to select actions based on neighboring states whenever the state has never been visited before. These are modifications that could further improve the agent.